

```

In [ ]: import numpy as np
import pandas as pd
import os
import pickle
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
import torchvision.datasets as datasets
from torch.utils.data import DataLoader, random_split, TensorDataset
from torch.optim.lr_scheduler import StepLR, MultiStepLR
from PIL import Image
import torch.optim.lr_scheduler as lr_scheduler
import matplotlib.pyplot as plt

# auto. choose CPU or GPU
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

# Function to load CIFAR-10 dataset
def load_cifar_batch(file):
    with open(file, 'rb') as fo:
        dict = pickle.load(fo, encoding='bytes')
    return dict

# Specify the directory containing CIFAR-10 batches
cifar10_dir = '/content/'
# Load metadata (labels)
meta_data_dict = load_cifar_batch(os.path.join(cifar10_dir, 'batches.meta'))
label_names = [label.decode('utf-8') for label in meta_data_dict[b'label_names']]

# Load training data
train_data = []
train_labels = []
for i in range(1, 6):
    batch = load_cifar_batch(os.path.join(cifar10_dir, f'data_batch_{i}'))
    train_data.append(batch[b'data'])
    train_labels += batch[b'labels']

train_data = np.vstack(train_data).reshape(-1, 3, 32, 32).transpose(0, 2, 3, 1) # Convert to HWC format
train_labels = np.array(train_labels)

# Data augmentation and normalization
transform = transforms.Compose([
    transforms.ToPILImage(), # Convert numpy array to PIL Image
    transforms.RandomRotation(10),
    transforms.ColorJitter(brightness = 0.1, contrast = 0.1, saturation = 0.1),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomAdjustSharpness(sharpness_factor = 2, p = 0.2),
    transforms.RandomCrop(32, padding=4),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.247, 0.243, 0.261)),
    transforms.RandomErasing(p=0.2, scale=(0.02, 0.1), value=1.0, inplace=False)
])

# Convert to TensorDataset and apply transformations
class CustomCIFAR10Dataset(torch.utils.data.Dataset):
    def __init__(self, images, labels, transform=None):
        self.images = images
        self.labels = labels
        self.transform = transform

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        img = self.images[idx]
        label = self.labels[idx]

        if self.transform:
            img = self.transform(img)

        return img, label

```

```

train_dataset = CustomCIFAR10Dataset(train_data, train_labels, transform=transform)

# Split into training and validation sets
#train_size = int(0.9 * len(train_dataset))
#val_size = len(train_dataset) - train_size
#train_dataset, val_dataset = random_split(train_dataset, [train_size, val_size])

test_transform = transforms.Compose([
    transforms.ToPILImage(), # Convert numpy array to PIL Image
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.247, 0.243, 0.261))
])

batch_test_dict = load_cifar_batch(os.path.join(cifar10_dir, 'test_batch'))
val_images = batch_test_dict[b'data'].reshape(-1, 3, 32, 32).transpose(0, 2, 3, 1)
val_labels = np.array(batch_test_dict[b'labels'])

val_dataset = CustomCIFAR10Dataset(val_images, val_labels, transform=test_transform)

# DataLoaders
train_loader = DataLoader(train_dataset, batch_size=256, shuffle=True, num_workers=4)
val_loader = DataLoader(val_dataset, batch_size=256, shuffle=False, num_workers=4)

# Load test dataset
cifar_test_path = '/content/cifar_test_nolabel.pkl'
test_batch = load_cifar_batch(cifar_test_path)
test_images = test_batch[b'data'].astype(np.float32) / 255.0

# Convert test dataset to Tensor
test_dataset = [(test_transform(img),) for img in test_images]
test_loader = DataLoader(test_dataset, batch_size=256, shuffle=False, num_workers=4)

# Train function + plot
def train_model(model, train_loader, val_loader, epochs=50):
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.9, weight_decay=1e-4)
    scheduler = MultiStepLR(optimizer, milestones=[30, 60, 80, 90], gamma=0.1)

    train_losses = []
    val_losses = []
    train_accuracies = []
    val_accuracies = []

    for epoch in range(epochs):
        # Training Phase
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0
        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        train_loss = running_loss / len(train_loader)
        train_acc = 100 * correct / total
        train_losses.append(train_loss)
        train_accuracies.append(train_acc)

        # Validation Phase
        model.eval()
        val_loss = 0.0
        correct = 0
        total = 0
        with torch.no_grad():
            for images, labels in val_loader:

```

```

        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)
        val_loss += loss.item()

        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    val_loss /= len(val_loader)
    val_acc = 100 * correct / total
    val_losses.append(val_loss)
    val_accuracies.append(val_acc)

    scheduler.step()
    print(f'Epoch {epoch+1}, Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.2f}%, Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%')

# Plot Losses
plt.figure(figsize=(10, 5))
plt.plot(range(1, epochs + 1), train_losses, label='Train Loss', color='red')
plt.plot(range(1, epochs + 1), val_losses, label='Validation Loss', color='blue')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Train Loss & Validation Loss')
plt.legend()
plt.grid()
plt.show()

# Plot Accuracies
plt.figure(figsize=(10, 5))
plt.plot(range(1, epochs + 1), train_accuracies, label='Train Accuracy', color='green')
plt.plot(range(1, epochs + 1), val_accuracies, label='Validation Accuracy', color='purple')
plt.xlabel('Epochs')
plt.ylabel('Accuracy (%)')
plt.title('Train Accuracy & Validation Accuracy')
plt.legend()
plt.grid()
plt.show()

import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
import torchvision.datasets as datasets
from torch.utils.data import DataLoader
import torch.optim.lr_scheduler as lr_scheduler

# Device configuration
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Define the Squeeze-and-Excitation (SE) Block
class SEBlock(nn.Module):
    def __init__(self, channels, reduction=16):
        super(SEBlock, self).__init__()
        self.global_avg_pool = nn.AdaptiveAvgPool2d(1)
        self.fc1 = nn.Linear(channels, channels // reduction, bias=False)
        self.relu = nn.ReLU(inplace=True)
        self.fc2 = nn.Linear(channels // reduction, channels, bias=False)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        b, c, _, _ = x.size()
        y = self.global_avg_pool(x).view(b, c)
        y = self.relu(self.fc1(y))
        y = self.sigmoid(self.fc2(y)).view(b, c, 1, 1)
        return x * y.expand_as(x)

# Define Residual Block with Shortcut Connections
class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_sizes, stride=1, use_se=True):
        super(ResidualBlock, self).__init__()
        mid_channels = out_channels // 2

        self.conv1 = nn.Conv2d(in_channels, mid_channels, kernel_size=kernel_sizes[0], padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(mid_channels)

```

```

self.conv2 = nn.Conv2d(mid_channels, mid_channels, kernel_size=kernel_sizes[1], padding=1, stride=1)
self.bn2 = nn.BatchNorm2d(mid_channels)

self.conv3 = nn.Conv2d(mid_channels, out_channels, kernel_size=kernel_sizes[2], padding=1, bias=False)
self.bn3 = nn.BatchNorm2d(out_channels)

self.se = SEBlock(out_channels) if use_se else nn.Identity()
self.leaky_relu = nn.LeakyReLU(0.1, inplace=True)
self.dropout = nn.Dropout2d(0.2)

self.shortcut = nn.Sequential()
if stride != 1 or in_channels != out_channels:
    self.shortcut = nn.Sequential(
        nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=stride, bias=False),
        nn.BatchNorm2d(out_channels)
    )

def forward(self, x):
    identity = self.shortcut(x)
    out = self.leaky_relu(self.bn1(self.conv1(x)))
    out = self.leaky_relu(self.bn2(self.conv2(out)))
    out = self.bn3(self.conv3(out))
    out = self.se(out)
    out = self.dropout(out)
    out += identity
    return self.leaky_relu(out)

class CustomResNet_v4(nn.Module):
    def __init__(self, num_classes=10):
        super(CustomResNet_v4, self).__init__()
        # Initial convolution layer with a small increase in channels
        self.init_conv = nn.Conv2d(3, 96, kernel_size=3, stride=1, padding=1, bias=False)
        self.init_bn = nn.BatchNorm2d(96)
        self.leaky_relu = nn.LeakyReLU(0.1, inplace=True)

        # First residual block with slightly increased channels
        self.layer1 = self._make_layer(96, 124, [3, 3, 3], 4, stride=1)
        # Second residual block with slightly increased channels
        self.layer2 = self._make_layer(124, 189, [3, 3, 3], 4, stride=2)
        # Third residual block with slightly increased channels
        self.layer3 = self._make_layer(189, 280, [3, 3, 3], 3, stride=2)

        # Final average pooling and fully connected layers
        self.avg_pool = nn.AdaptiveAvgPool2d(1)
        self.fc = nn.Linear(280, num_classes)

    def _make_layer(self, in_channels, out_channels, kernel_sizes, blocks, stride):
        layers = [ResidualBlock(in_channels, out_channels, kernel_sizes, stride)]
        for _ in range(1, blocks):
            layers.append(ResidualBlock(out_channels, out_channels, kernel_sizes, stride=1))
        return nn.Sequential(*layers)

    def forward(self, x):
        out = self.leaky_relu(self.init_bn(self.init_conv(x)))
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.avg_pool(out)
        out = torch.flatten(out, 1)
        out = self.fc(out)
        return out

# Instantiate the model
model = CustomResNet_v4().to(device)

# Define the optimizer, scheduler, and loss function
optimizer = optim.SGD(model.parameters(), lr=0.05, momentum=0.9, weight_decay=0.0005)
scheduler = lr_scheduler.CosineAnnealingLR(optimizer, T_max=100)
criterion = nn.CrossEntropyLoss()

# Print the model summary
from torchsummary import summary

```

```
summary(model, (3, 32, 32))

# Train the model
train_model(model, train_loader, val_loader, epochs=70) #change epoch

# Generate submission file
model.eval()
predictions = []
with torch.no_grad():
    for batch in test_loader:
        images = batch[0].to(device) # Get images tensor from tuple and move to device
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        predictions.extend(predicted.cpu().numpy())

# Generate submission file
submission = pd.DataFrame({'ID': np.arange(len(predictions)), 'Labels': predictions})
submission.to_csv('/content/submission1.csv', index=False)
print("Submission1 file saved.")
```

Using device: cuda

/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

warnings.warn(

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 96, 32, 32]	2,592
BatchNorm2d-2	[-1, 96, 32, 32]	192
LeakyReLU-3	[-1, 96, 32, 32]	0
Conv2d-4	[-1, 124, 32, 32]	11,904
BatchNorm2d-5	[-1, 124, 32, 32]	248
Conv2d-6	[-1, 62, 32, 32]	53,568
BatchNorm2d-7	[-1, 62, 32, 32]	124
LeakyReLU-8	[-1, 62, 32, 32]	0
Conv2d-9	[-1, 62, 32, 32]	34,596
BatchNorm2d-10	[-1, 62, 32, 32]	124
LeakyReLU-11	[-1, 62, 32, 32]	0
Conv2d-12	[-1, 124, 32, 32]	69,192
BatchNorm2d-13	[-1, 124, 32, 32]	248
AdaptiveAvgPool2d-14	[-1, 124, 1, 1]	0
Linear-15	[-1, 7]	868
ReLU-16	[-1, 7]	0
Linear-17	[-1, 124]	868
Sigmoid-18	[-1, 124]	0
SEBlock-19	[-1, 124, 32, 32]	0
Dropout2d-20	[-1, 124, 32, 32]	0
LeakyReLU-21	[-1, 124, 32, 32]	0
ResidualBlock-22	[-1, 124, 32, 32]	0
Conv2d-23	[-1, 62, 32, 32]	69,192
BatchNorm2d-24	[-1, 62, 32, 32]	124
LeakyReLU-25	[-1, 62, 32, 32]	0
Conv2d-26	[-1, 62, 32, 32]	34,596
BatchNorm2d-27	[-1, 62, 32, 32]	124
LeakyReLU-28	[-1, 62, 32, 32]	0
Conv2d-29	[-1, 124, 32, 32]	69,192
BatchNorm2d-30	[-1, 124, 32, 32]	248
AdaptiveAvgPool2d-31	[-1, 124, 1, 1]	0
Linear-32	[-1, 7]	868
ReLU-33	[-1, 7]	0
Linear-34	[-1, 124]	868
Sigmoid-35	[-1, 124]	0
SEBlock-36	[-1, 124, 32, 32]	0
Dropout2d-37	[-1, 124, 32, 32]	0
LeakyReLU-38	[-1, 124, 32, 32]	0
ResidualBlock-39	[-1, 124, 32, 32]	0
Conv2d-40	[-1, 62, 32, 32]	69,192
BatchNorm2d-41	[-1, 62, 32, 32]	124
LeakyReLU-42	[-1, 62, 32, 32]	0
Conv2d-43	[-1, 62, 32, 32]	34,596
BatchNorm2d-44	[-1, 62, 32, 32]	124
LeakyReLU-45	[-1, 62, 32, 32]	0
Conv2d-46	[-1, 124, 32, 32]	69,192
BatchNorm2d-47	[-1, 124, 32, 32]	248
AdaptiveAvgPool2d-48	[-1, 124, 1, 1]	0
Linear-49	[-1, 7]	868
ReLU-50	[-1, 7]	0
Linear-51	[-1, 124]	868
Sigmoid-52	[-1, 124]	0
SEBlock-53	[-1, 124, 32, 32]	0
Dropout2d-54	[-1, 124, 32, 32]	0
LeakyReLU-55	[-1, 124, 32, 32]	0
ResidualBlock-56	[-1, 124, 32, 32]	0
Conv2d-57	[-1, 62, 32, 32]	69,192
BatchNorm2d-58	[-1, 62, 32, 32]	124
LeakyReLU-59	[-1, 62, 32, 32]	0
Conv2d-60	[-1, 62, 32, 32]	34,596
BatchNorm2d-61	[-1, 62, 32, 32]	124
LeakyReLU-62	[-1, 62, 32, 32]	0
Conv2d-63	[-1, 124, 32, 32]	69,192
BatchNorm2d-64	[-1, 124, 32, 32]	248
AdaptiveAvgPool2d-65	[-1, 124, 1, 1]	0
Linear-66	[-1, 7]	868
ReLU-67	[-1, 7]	0
Linear-68	[-1, 124]	868
Sigmoid-69	[-1, 124]	0
SEBlock-70	[-1, 124, 32, 32]	0
Dropout2d-71	[-1, 124, 32, 32]	0
LeakyReLU-72	[-1, 124, 32, 32]	0

ResidualBlock-73	[-1, 124, 32, 32]	0
Conv2d-74	[-1, 189, 16, 16]	23,436
BatchNorm2d-75	[-1, 189, 16, 16]	378
Conv2d-76	[-1, 94, 32, 32]	104,904
BatchNorm2d-77	[-1, 94, 32, 32]	188
LeakyReLU-78	[-1, 94, 32, 32]	0
Conv2d-79	[-1, 94, 16, 16]	79,524
BatchNorm2d-80	[-1, 94, 16, 16]	188
LeakyReLU-81	[-1, 94, 16, 16]	0
Conv2d-82	[-1, 189, 16, 16]	159,894
BatchNorm2d-83	[-1, 189, 16, 16]	378
AdaptiveAvgPool2d-84	[-1, 189, 1, 1]	0
Linear-85	[-1, 11]	2,079
ReLU-86	[-1, 11]	0
Linear-87	[-1, 189]	2,079
Sigmoid-88	[-1, 189]	0
SEBlock-89	[-1, 189, 16, 16]	0
Dropout2d-90	[-1, 189, 16, 16]	0
LeakyReLU-91	[-1, 189, 16, 16]	0
ResidualBlock-92	[-1, 189, 16, 16]	0
Conv2d-93	[-1, 94, 16, 16]	159,894
BatchNorm2d-94	[-1, 94, 16, 16]	188
LeakyReLU-95	[-1, 94, 16, 16]	0
Conv2d-96	[-1, 94, 16, 16]	79,524
BatchNorm2d-97	[-1, 94, 16, 16]	188
LeakyReLU-98	[-1, 94, 16, 16]	0
Conv2d-99	[-1, 189, 16, 16]	159,894
BatchNorm2d-100	[-1, 189, 16, 16]	378
AdaptiveAvgPool2d-101	[-1, 189, 1, 1]	0
Linear-102	[-1, 11]	2,079
ReLU-103	[-1, 11]	0
Linear-104	[-1, 189]	2,079
Sigmoid-105	[-1, 189]	0
SEBlock-106	[-1, 189, 16, 16]	0
Dropout2d-107	[-1, 189, 16, 16]	0
LeakyReLU-108	[-1, 189, 16, 16]	0
ResidualBlock-109	[-1, 189, 16, 16]	0
Conv2d-110	[-1, 94, 16, 16]	159,894
BatchNorm2d-111	[-1, 94, 16, 16]	188
LeakyReLU-112	[-1, 94, 16, 16]	0
Conv2d-113	[-1, 94, 16, 16]	79,524
BatchNorm2d-114	[-1, 94, 16, 16]	188
LeakyReLU-115	[-1, 94, 16, 16]	0
Conv2d-116	[-1, 189, 16, 16]	159,894
BatchNorm2d-117	[-1, 189, 16, 16]	378
AdaptiveAvgPool2d-118	[-1, 189, 1, 1]	0
Linear-119	[-1, 11]	2,079
ReLU-120	[-1, 11]	0
Linear-121	[-1, 189]	2,079
Sigmoid-122	[-1, 189]	0
SEBlock-123	[-1, 189, 16, 16]	0
Dropout2d-124	[-1, 189, 16, 16]	0
LeakyReLU-125	[-1, 189, 16, 16]	0
ResidualBlock-126	[-1, 189, 16, 16]	0
Conv2d-127	[-1, 94, 16, 16]	159,894
BatchNorm2d-128	[-1, 94, 16, 16]	188
LeakyReLU-129	[-1, 94, 16, 16]	0
Conv2d-130	[-1, 94, 16, 16]	79,524
BatchNorm2d-131	[-1, 94, 16, 16]	188
LeakyReLU-132	[-1, 94, 16, 16]	0
Conv2d-133	[-1, 189, 16, 16]	159,894
BatchNorm2d-134	[-1, 189, 16, 16]	378
AdaptiveAvgPool2d-135	[-1, 189, 1, 1]	0
Linear-136	[-1, 11]	2,079
ReLU-137	[-1, 11]	0
Linear-138	[-1, 189]	2,079
Sigmoid-139	[-1, 189]	0
SEBlock-140	[-1, 189, 16, 16]	0
Dropout2d-141	[-1, 189, 16, 16]	0
LeakyReLU-142	[-1, 189, 16, 16]	0
ResidualBlock-143	[-1, 189, 16, 16]	0
Conv2d-144	[-1, 280, 8, 8]	52,920
BatchNorm2d-145	[-1, 280, 8, 8]	560
Conv2d-146	[-1, 140, 16, 16]	238,140
BatchNorm2d-147	[-1, 140, 16, 16]	280

LeakyReLU-148	[-1, 140, 16, 16]	0
Conv2d-149	[-1, 140, 8, 8]	176,400
BatchNorm2d-150	[-1, 140, 8, 8]	280
LeakyReLU-151	[-1, 140, 8, 8]	0
Conv2d-152	[-1, 280, 8, 8]	352,800
BatchNorm2d-153	[-1, 280, 8, 8]	560
AdaptiveAvgPool2d-154	[-1, 280, 1, 1]	0
Linear-155	[-1, 17]	4,760
ReLU-156	[-1, 17]	0
Linear-157	[-1, 280]	4,760
Sigmoid-158	[-1, 280]	0
SEBlock-159	[-1, 280, 8, 8]	0
Dropout2d-160	[-1, 280, 8, 8]	0
LeakyReLU-161	[-1, 280, 8, 8]	0
ResidualBlock-162	[-1, 280, 8, 8]	0
Conv2d-163	[-1, 140, 8, 8]	352,800
BatchNorm2d-164	[-1, 140, 8, 8]	280
LeakyReLU-165	[-1, 140, 8, 8]	0
Conv2d-166	[-1, 140, 8, 8]	176,400
BatchNorm2d-167	[-1, 140, 8, 8]	280
LeakyReLU-168	[-1, 140, 8, 8]	0
Conv2d-169	[-1, 280, 8, 8]	352,800
BatchNorm2d-170	[-1, 280, 8, 8]	560
AdaptiveAvgPool2d-171	[-1, 280, 1, 1]	0
Linear-172	[-1, 17]	4,760
ReLU-173	[-1, 17]	0
Linear-174	[-1, 280]	4,760
Sigmoid-175	[-1, 280]	0
SEBlock-176	[-1, 280, 8, 8]	0
Dropout2d-177	[-1, 280, 8, 8]	0
LeakyReLU-178	[-1, 280, 8, 8]	0
ResidualBlock-179	[-1, 280, 8, 8]	0
Conv2d-180	[-1, 140, 8, 8]	352,800
BatchNorm2d-181	[-1, 140, 8, 8]	280
LeakyReLU-182	[-1, 140, 8, 8]	0
Conv2d-183	[-1, 140, 8, 8]	176,400
BatchNorm2d-184	[-1, 140, 8, 8]	280
LeakyReLU-185	[-1, 140, 8, 8]	0
Conv2d-186	[-1, 280, 8, 8]	352,800
BatchNorm2d-187	[-1, 280, 8, 8]	560
AdaptiveAvgPool2d-188	[-1, 280, 1, 1]	0
Linear-189	[-1, 17]	4,760
ReLU-190	[-1, 17]	0
Linear-191	[-1, 280]	4,760
Sigmoid-192	[-1, 280]	0
SEBlock-193	[-1, 280, 8, 8]	0
Dropout2d-194	[-1, 280, 8, 8]	0
LeakyReLU-195	[-1, 280, 8, 8]	0
ResidualBlock-196	[-1, 280, 8, 8]	0
AdaptiveAvgPool2d-197	[-1, 280, 1, 1]	0
Linear-198	[-1, 10]	2,810

=====

Total params: 4,905,430

Trainable params: 4,905,430

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 59.35

Params size (MB): 18.71

Estimated Total Size (MB): 78.08

Epoch 1, Train Loss: 1.7110, Train Acc: 35.65%, Val Loss: 1.5430, Val Acc: 42.16%

Epoch 2, Train Loss: 1.3212, Train Acc: 51.65%, Val Loss: 1.2077, Val Acc: 57.29%

Epoch 3, Train Loss: 1.1145, Train Acc: 59.97%, Val Loss: 0.9759, Val Acc: 65.78%

Epoch 4, Train Loss: 0.9789, Train Acc: 65.34%, Val Loss: 0.9797, Val Acc: 66.82%

Epoch 5, Train Loss: 0.8809, Train Acc: 68.88%, Val Loss: 0.9081, Val Acc: 68.96%

Epoch 6, Train Loss: 0.7801, Train Acc: 72.53%, Val Loss: 0.8567, Val Acc: 72.06%

Epoch 7, Train Loss: 0.7115, Train Acc: 75.11%, Val Loss: 0.7762, Val Acc: 74.21%

Epoch 8, Train Loss: 0.6495, Train Acc: 77.41%, Val Loss: 0.5807, Val Acc: 80.84%

Epoch 9, Train Loss: 0.5967, Train Acc: 79.02%, Val Loss: 0.5482, Val Acc: 80.98%

Epoch 10, Train Loss: 0.5628, Train Acc: 80.35%, Val Loss: 0.4909, Val Acc: 83.77%



Epoch 11, Train Loss: 0.5284, Train Acc: 81.51%, Val Loss: 0.5318, Val Acc: 82.39%

Epoch 12, Train Loss: 0.5024, Train Acc: 82.61%, Val Loss: 0.6906, Val Acc: 78.11%

Epoch 13, Train Loss: 0.4754, Train Acc: 83.40%, Val Loss: 0.4482, Val Acc: 84.63%

Epoch 14, Train Loss: 0.4463, Train Acc: 84.53%, Val Loss: 0.4074, Val Acc: 86.25%

Epoch	Train Loss	Train Acc	Val Loss	Val Acc
Epoch 15,	0.4351,	84.91%,	0.4760,	83.79%
Epoch 16,	0.4241,	85.33%,	0.4498,	84.83%
Epoch 17,	0.4039,	85.97%,	0.3965,	86.52%
Epoch 18,	0.3816,	86.90%,	0.4165,	86.04%
Epoch 19,	0.3717,	86.98%,	0.4206,	86.48%
Epoch 20,	0.3632,	87.49%,	0.3898,	87.13%
Epoch 21,	0.3446,	88.10%,	0.3817,	87.38%
Epoch 22,	0.3346,	88.34%,	0.3460,	88.61%
Epoch 23,	0.3257,	88.60%,	0.3226,	89.50%
Epoch 24,	0.3162,	88.95%,	0.3312,	88.74%
Epoch 25,	0.3107,	89.24%,	0.3553,	88.75%
Epoch 26,	0.2937,	89.80%,	0.3214,	89.33%
Epoch 27,	0.2924,	89.76%,	0.3529,	88.84%
Epoch 28,	0.2816,	90.11%,	0.3646,	88.34%
Epoch 29,	0.2769,	90.33%,	0.3054,	90.07%
Epoch 30,	0.2715,	90.49%,	0.2846,	90.40%
Epoch 31,	0.1977,	93.12%,	0.2199,	92.52%
Epoch 32,	0.1700,	94.12%,	0.2126,	92.68%
Epoch 33,	0.1634,	94.35%,	0.2135,	92.86%
Epoch 34,	0.1576,	94.59%,	0.2118,	92.76%
Epoch 35,	0.1499,	94.76%,	0.2098,	92.97%
Epoch 36,	0.1470,	94.97%,	0.2116,	92.97%
Epoch 37,	0.1400,	95.14%,	0.2136,	93.16%
Epoch 38,	0.1399,	95.07%,	0.2078,	93.22%
Epoch 39,	0.1383,	95.14%,	0.2088,	93.11%
Epoch 40,	0.1349,	95.22%,	0.2041,	93.31%
Epoch 41,	0.1311,	95.44%,	0.2088,	93.18%
Epoch 42,	0.1286,	95.54%,	0.2056,	93.18%
Epoch 43,	0.1238,	95.66%,	0.2052,	93.31%
Epoch 44,	0.1232,	95.69%,	0.2104,	93.36%
Epoch 45,	0.1181,	95.96%,	0.2100,	93.30%
Epoch 46,	0.1183,	95.86%,	0.2078,	93.27%
Epoch 47,	0.1178,	95.93%,	0.2085,	93.18%
Epoch 48,	0.1151,	95.92%,	0.2080,	93.43%
Epoch 49,	0.1144,	96.03%,	0.2192,	93.14%
Epoch 50,	0.1133,	96.03%,	0.2029,	93.54%
Epoch 51,	0.1102,	96.14%,	0.2059,	93.64%
Epoch 52,	0.1065,	96.19%,	0.2083,	93.43%
Epoch 53,	0.1071,	96.21%,	0.2071,	93.58%
Epoch 54,	0.1075,	96.24%,	0.2114,	93.40%
Epoch 55,	0.1038,	96.44%,	0.2093,	93.74%
Epoch 56,	0.1023,	96.47%,	0.2054,	93.67%
Epoch 57,	0.1027,	96.43%,	0.2083,	93.38%
Epoch 58,	0.0991,	96.58%,	0.2069,	93.56%
Epoch 59,	0.0980,	96.50%,	0.2124,	93.61%
Epoch 60,	0.0970,	96.51%,	0.2098,	93.65%
Epoch 61,	0.0888,	96.93%,	0.2050,	93.75%
Epoch 62,	0.0881,	96.89%,	0.2032,	93.65%
Epoch 63,	0.0867,	97.06%,	0.2037,	93.73%
Epoch 64,	0.0865,	97.03%,	0.2000,	93.81%
Epoch 65,	0.0827,	97.14%,	0.2023,	93.77%
Epoch 66,	0.0808,	97.21%,	0.2013,	93.83%
Epoch 67,	0.0821,	97.17%,	0.2023,	93.88%
Epoch 68,	0.0807,	97.24%,	0.2005,	93.83%
Epoch 69,	0.0818,	97.19%,	0.2057,	93.93%
Epoch 70,	0.0803,	97.25%,	0.2019,	93.98%



Submission1 file saved.

```
In [ ]: torch.save(model.state_dict(), '/content/model_checkpoint.pth')
```

```
In [ ]: # Generate submission file
model.eval()
predictions = []
with torch.no_grad():
    for batch in test_loader:
        images = batch[0].to(device) # Get images tensor from tuple and move to device
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        predictions.extend(predicted.cpu().numpy())

# Generate submission file
submission = pd.DataFrame({'ID': np.arange(len(predictions)), 'Labels': predictions})
submission.to_csv('/content/submission2.csv', index=False)
print("Submission1 file saved.")
```

Submission1 file saved.

In []: